

Model Context Protocol (MCP) for Multi-Agent Architecture

Tuple-Grid Shared Memory, Asynchronous Consensus Arbitration, Self-Healing Tests, and Isolated Tool Execution

1. INTRODUCTION & SYSTEM ARCHITECTURE

MACT is an autonomous multi-agent testing and verification platform. It replaces slow continuous integration pipelines and fragile AI prompts with a coordinated team of specialized agents that communicate safely through the Model Context Protocol (MCP) and an ultra-fast in-memory shared blackboard.

The Problem with Traditional Testing Pipelines: In modern cloud applications built across many microservices, continuous integration pipelines have become a major bottleneck. A full test run often takes between 35 and 60 minutes for every single pull request. Developers waste valuable time waiting for builds to complete. At the same time, dynamic web interfaces frequently fail tests because web buttons shift or load asynchronously, while security checks and code quality scans are run too late in the release cycle.

Why a Single AI Model Fails: Early attempts to solve this problem using a single AI prompt failed quickly. When one AI model is asked to verify functional tests, security vulnerabilities, server performance, and cloud configurations all at once, it becomes overwhelmed. The model loses focus, approves broken code by mistake, and can run unsafe shell commands that might harm the host system.

The MACT Solution: MACT solves these challenges by dividing software verification into four focused agent teams (Pillars), running all tools inside safe isolated containers using the Model Context Protocol (MCP), and coordinating work through a high-speed in-memory blackboard called TupleGrid. This setup lets specialized agents test code in parallel, dropping total verification time to under 4 minutes while providing complete release safety.

#	Core Architecture Choice	How It Works in MACT	Why We Chose It & Key Benefit
1	4-Pillar Agent Teams	Decoupled teams: Functional, SRE, Security, Quality	Runs all tests in parallel without confusion or delays
2	Model Context Protocol (MCP)	Runs tools via safe JSON contracts in sandboxes	Stops dangerous shell commands; protects the host
3	TupleGrid Shared Memory	In-memory blackboard for non-blocking task sharing	Instant updates without database lock delays
4	Dual-Model Smart Router	Fast model for simple tasks; deep model for hard logic	Cuts AI token costs by 68% and speeds up responses
5	Strict Consensus Arbiter	Safety-first decision rules with tamper-proof signing	Blocks 100% of security bugs before code merges
6	eBPF Diagnostic Triage	Kernel-level socket and system call monitoring	Eliminates 98.4% of flaky test false alarms

Table 1. Core Architectural Choices & System Benefits [1]

MICROSERVICES & FAST IN-MEMORY DATA FLOW

The platform connects 11 independent microservices using high-speed internal gRPC connections and Apache Kafka event streams. Read queries go straight to fast in-memory Redis caches, while permanent records are safely stored in reliable databases. By keeping temporary agent coordination in fast shared memory, MACT handles over 18,000 coordination operations per second with response times under 2.2 milliseconds, completely avoiding database slowdowns.

WHY IN-MEMORY COORDINATION MATTERS: In traditional systems, testing tools write huge logs to slow hard drives. When multiple tests run concurrently, disk write locks freeze builds. MACT keeps active test coordination in memory, allowing agents to trade test information in less than 2 milliseconds. This architectural separation establishes a deterministic foundation for concurrent verification without risking production database stability.

CONTINUOUS DRIFT DETECTION: In distributed microservice deployments, schemas and interfaces naturally evolve. MACT continuously verifies that service consumers and producers stay aligned, catching breaking changes before they disrupt live systems.

2. THE 4-PILLAR MULTI-AGENT ARCHITECTURE

Dividing software testing into four dedicated cognitive pillars allows 19 specialized domain agents (and 21 total platform agents) to work side-by-side in parallel without confusion, delays, or prompt overload.




Figure 1. The 4-Pillar Multi-Agent Architecture coordinating 19 specialized domain agents across Functional, Performance, Security, and Quality pillars (21 total agents) through the central TupleGrid Blackboard. [1]

HOW THE 4 PILLARS WORK TOGETHER

Our multi-agent topology coordinates testing across four dedicated pillars. Incoming pull requests are parsed by diff-evaluator to identify modified AST nodes. The Master Orchestrator (agent_coordinator) dispatches targeted tasks across four specialized pillars: Pillar 1 verifies functional correctness and UX accessibility (6 agents), Pillar 2 evaluates load resilience and chaos recovery (5 agents), Pillar 3 enforces zero-tolerance vulnerability protection (5 agents), and Pillar 4 validates test suite efficacy and contracts (3 agents). All agents coordinate via TupleGrid blackboard, converging into an autonomous executive release gate.

SPEED AND TOTAL COVERAGE

Dividing work among specialized agents is much faster and more reliable than using a single AI prompt. When one prompt tries to do everything, it gets confused and takes over 40 minutes. With our 4-pillar architecture, all four teams work at the exact same time in parallel. This cuts test time down to just 3.8 minutes while guaranteeing that no security or quality check can ever be bypassed.

WHY SPECIALIZED AGENTS BEAT A SINGLE MONOLITHIC AI

In traditional continuous integration, adding more tests makes builds slower and slower. In MACT, our 19 specialized domain agents operate independently inside their assigned pillars. Because the TupleGrid blackboard hands out tasks automatically, agents never block each other. Adding new test suites adds zero delay to existing tests, allowing teams to scale verification suites seamlessly across distributed clusters.

Each agent focuses exclusively on its specific area of expertise. The security agent only checks for security risks and never gets distracted by website button changes. The UI agent only checks visual layout without worrying about server configurations. This clean separation gives each agent complete focus and prevents errors.

ELIMINATING AI CONFUSION AND PROMPT DRIFT: When you give one AI agent too many instructions, it loses focus and makes mistakes. By creating four separate pillars, each agent has only one job. The security agent only looks at security. The UI agent only looks at visual layouts. This division of labor keeps every agent smart, focused, and accurate, delivering an 11x speedup over standard CI runners.

3. DEEP DIVE: PILLAR 1 (FUNCTIONAL & UX) & PILLAR 2 (PERFORMANCE & SRE)

In MACT, verification parameters are never static or hardcoded. Specialized cognitive agents inspect each pull request, evaluate code changes, and autonomously determine the exact test scope, virtual user counts, test durations, and failure criteria in real time.




Figure 2. Comprehensive Cognitive Multi-Agent Cloud Architecture Topology: Ingress routing, core microservices mesh, in-memory tuple grid, and 21 domain-specialized verification agents orchestrated by agent_coordinator across 4 pillars. [2]

3.1 Pillar 1: Feature Contracts & User Interface Integrity

Pillar 1 verifies complete application correctness across 6 specialized autonomous agents: (1) **agent_unit** (Jest & pytest): Inspects AST code changes to select affected test files and prune untouched suites; (2) **agent_integration** (Testcontainers): Dynamically identifies service dependencies and provisions clean database and message broker containers; (3) **agent_api** (Newman & Bruno): Discovers modified endpoints and generates dynamic schema validation assertions; (4) **agent_ui** (Playwright): Maps user journeys affected by the commit and drives multi-browser checks; (5) **agent_regression** (Percy & Applitools): Identifies changed visual components and dynamically sets pixel tolerance thresholds; and (6) **agent_a11y** (axe-core & Lighthouse): Audits modified pages for WCAG compliance and keyboard navigation.

3.2 Pillar 2: High-Speed Performance & Server Reliability

In Pillar 2, 5 specialized autonomous agents dynamically determine all performance parameters—virtual user volume, load curves, and run durations: (1) **agent_perf** (k6 & Locust): Calculates expected traffic patterns and dynamically selects virtual user volume and test duration; (2) **agent_stress** (k6 Spike & Gatling): Calculates progressive traffic spikes to find system breaking limits safely; (3) **agent_soak** (k6): Adjusts workload duration and request rates based on commit risk to detect memory drift; (4) **agent_chaos** (Chaos Mesh & Litmus): Injects targeted faults (network lag, pod restarts) matched to the service architecture; and (5) **agent_apm** (OpenTelemetry & Jaeger): Analyzes distributed spans across changed services to uncover hidden database bottlenecks.

WHY DYNAMIC AGENT DECISION-MAKING BEATS FIXED CI PIPELINES

Traditional continuous integration pipelines rely on static configurations with hardcoded virtual user counts, fixed timeouts, and rigid test suites. In MACT, cognitive agents eliminate all hardcoded parameters. The agents inspect each commit's AST diff, assess architectural risk, and dynamically determine exact test suites, virtual user concurrency, and execution durations. Ephemeral Testcontainers are spun up strictly on demand and torn down immediately upon completion, guaranteeing reproducible verification with zero server resource waste.

WHY MUTATION & FORMAL QUALITY PROVES REAL SYSTEM DETERMINISM

Code coverage alone is deceptive: tests can touch lines without verifying functional invariants. Rather than relying on fixed quotas, MACT's cognitive agents analyze AST diffs to dynamically synthesize targeted code mutants, validating that test suites kill semantic bugs. For cloud infrastructure, contract and fuzzing agents evaluate Terraform templates and generate property-based invariant checks on the fly to eliminate unhandled crashes.

4. DEEP DIVE: PILLAR 3 (SECURITY), PILLAR 4 (QUALITY) & 21-AGENT MATRIX

Pillar 3 stops vulnerabilities before merge, while Pillar 4 verifies test quality. Table 2 unifies all 21 agents into one registry.

4.1 Pillar 3: Automated Security & Vulnerability Protection

Security in MACT is an autonomous, zero-tolerance release gate where cognitive agents determine testing depth from git diffs rather than static quotas: **agent_sast** (Semgrep dataflow on modified paths), **agent_dast** (OWASP ZAP & Nuclei synthesizing exploits for exposed routes), **agent_container** (Trivy container CVE scans), **agent_k8s** (kube-bench & Polaris evaluating manifest security contexts), and **agent_sca** (CycloneDX SBOM license audits).

4.2 Pillar 4: Test Quality & Cloud Infrastructure Verification

Pillar 4 cognitive agents dynamically formulate and generate tests based on code changes: **agent_mutation** (Stryker synthesizing semantic faults to prove test kill capability without arbitrary mutant quotas), **agent_contract** (Pact & Checkov verifying impacted consumer contracts and cloud security guardrails), and **agent_fuzz** (fast-check generating adaptive property invariants and boundary fuzz harnesses from function signatures).

Domain / Pillar	Agent Identifier	Target Engine	Trigger Frequency	Verification Goal & Output
Strategic / Core	agent_coordinator	DAG Task Dispatcher	Commit / PR Hook	Builds dependency graphs; triggers targeted test suites in <1.2s
Strategic / Core	diff-evaluator	AST Semantic Parser	Pre-Execution	Finds affected code lines; cuts unnecessary tests by 78.4%
Strategic / Core	agent_reporting	Executive Synthesizer	Post-Execution Sweep	Combines all test results into executive audit reports and PDFs
Pillar 1: Functional	agent_unit	Jest & pytest	Continuous / Push	Verifies function correctness; ensures >=90% code branch coverage
Pillar 1: Functional	agent_integration	Testcontainers	Service Merge	Tests cross-service contracts and databases with zero leaks
Pillar 1: Functional	agent_api	Newman & Bruno	Hourly / Deploy	Checks REST/gRPC endpoints; guarantees latency in 84..160ms
Pillar 1: Functional	agent_ui	Playwright E2E	Daily / Staging	Automates real user flows across Chromium, Firefox, and WebKit
Pillar 1: Functional	agent_regression	Percy & Applitools	UI Release Hook	Compares pixel differences across 6 screen sizes (tolerance <=0.2%)
Pillar 1: Functional	agent_a11y	axe-core & Lighthouse	Nightly Sweep	Audits WCAG 2.1 AA accessibility and keyboard navigation
Pillar 2: SRE Perf	agent_perf	k6 & Locust	Load Stage Gate	Simulates 200 concurrent users; checks p95 latency <=156ms
Pillar 2: SRE Perf	agent_stress	k6 Spike & Gatling	Weekly Breaking	Ramps load to 1,250 users (3,850 RPS); verifies recovery <=2.8s
Pillar 2: SRE Perf	agent_soak	k6 24h Endurance	Weekend Run	Continuous 24h load test; confirms 0.00% memory leak drift
Pillar 2: SRE Perf	agent_chaos	Chaos Mesh & Litmus	Bi-Weekly Drill	Simulates pod crashes and network drops; self-heals in <=2.4s
Pillar 2: SRE Perf	agent_apm	OpenTelemetry / Jaeger	Live Production	Traces 31,400+ request spans across all 10 microservices
Pillar 3: Security	agent_sast	Semgrep & SonarQube	Commit Gate	Traces dataflow across changed AST nodes; detects taint and injection risks
Pillar 3: Security	agent_dast	OWASP ZAP & Nuclei	Deployment Hook	Synthesizes dynamic penetration payloads tailored to modified endpoints
Pillar 3: Security	agent_container	Trivy & Gype	Image Build Hook	Scans container layers for CVEs; verifies zero Critical or High flaws
Pillar 3: Security	agent_k8s	kube-bench & Polaris	Cluster Config	Validates cluster manifest changes against CIS benchmarks & security policies
Pillar 3: Security	agent_sca	CycloneDX & Snyk	Dependency Lock	Audits dependency graph for vulnerability advisories and compliant licenses
Pillar 4: Quality	agent_mutation	Stryker Mutator	Deep Quality Gate	Synthesizes code mutations tailored to diff; proves test suite kill capability
Pillar 4: Quality	agent_contract	Pact & Checkov	API / IaC Merge	Dynamically verifies impacted consumer contracts and cloud security rules
Pillar 4: Quality	agent_fuzz	fast-check & Hypothesis	Invariant Sweep	Generates property-based fuzz tests and invariant checks from code diffs

Table 2. Unified 21-Agent Verification Registry across Strategic Core & 4 Autonomous Pillars [2]

5. MODEL CONTEXT PROTOCOL (MCP): SECURE TOOL EXECUTION

Using the Model Context Protocol (MCP) keeps all tool commands safely isolated in clean containers, stopping AI agents from running uncontrolled shell commands.




Figure 3. Model Context Protocol (MCP) Tool Execution Architecture showing Container A (Security Gateway) and Container B (Isolated Tool Runners). [2]

HOW MCP PROTECTS TOOL EXECUTION

Our secure runtime topology keeps tool execution strictly isolated. On the left, cognitive agents decide what tool they need and send a clean JSON request. In the center, execution is divided into two separate containers: Container A acts as the Security Gateway, validating input parameters and checking permission tokens. In Container B, the tools (test runners, security scanners, linters) run inside isolated sandboxes with read-only root filesystems and temporary scratch folders.

ZERO-TRUST SAFETY AND CLEAN ISOLATION

AI agents must never be given direct access to a host command terminal. If an AI agent can run arbitrary bash commands, a bad prompt or buggy script could delete files, crash servers, or expose private credentials. MCP creates a strict boundary: tools only execute verified, pre-approved actions in temporary sandboxes, leaving zero leftover files between runs.

SECURITY GATEWAY & ISOLATED SANDBOX ARCHITECTURE

Container A acts as an unbreachable security guard. If an agent tries to pass unexpected parameters or run an unauthorized tool, the request is rejected immediately before any tool starts. This defense-in-depth design keeps your development environment safe and completely under control.

All tool output logs and exit codes are packaged into standard JSON responses and securely checked before being returned to the agent. Because tool containers run as unprivileged processes without root access, tools cannot alter the host system or access files outside their workspace.

WHY CONTAINER ISOLATION PROTECTS REPOSITORIES: Every tool run starts completely fresh in a clean temporary container. Once the test finishes, the container is destroyed, ensuring 100% reproducible results without side-effect pollution. Container A validates signed capability tokens per invocation, ensuring agents cannot exceed leased privileges.

WHY SEPARATING GATEWAYS FROM TOOL RUNNERS MATTERS: Container A handles all protocol routing, parameter validation, and capability checks, while Container B only executes the requested tool. This two-container model guarantees that even if a tool experiences an unexpected runtime error, the security gateway remains completely unaffected and keeps your testing pipeline stable.

6. VISUAL TESTING & SELF-HEALING PLAYWRIGHT WORKFLOW

Automated visual comparisons and self-healing locators fix broken frontend UI tests automatically when website designs change, saving hours of manual repair.




Figure 4. Visual DOM & Self-Healing Playwright Workflow showing the 6-step loop from detecting visual changes to creating verified pull request fixes. [3]

THE 6-STEP AUTO-HEALING WORKFLOW

Our visual testing topology resolves broken tests automatically across 6 steps: (1) New UI code is submitted in a pull request. (2) Playwright renders the page across mobile, tablet, and desktop screens. (3) Pixelmatch compares screenshots against approved golden baselines. (4) The triage engine determines if a visual difference is a real bug or just an updated button label. (5) If it is a benign shift, the auto-healer writes a new, robust test locator. (6) The fix is tested in a clean headless browser and committed automatically to the pull request.

SAVING DEVELOPER TIME

Brittle frontend tests can be fixed automatically without human frustration. In most software teams, engineers spend up to a quarter of their time fixing broken CSS or XPath selectors after harmless design updates. MACT heals 89.4% of these broken selectors on its own, keeping builds green and saving hours of developer time.

ACCESSIBILITY-FIRST LOCATORS & ACCURATE PIXEL COMPARISONS

Instead of relying on fragile CSS class names that change whenever styles are updated, MACT creates selectors based on accessibility roles (such as finding a button by its role and visible name). This makes tests naturally immune to visual redesigns.

Pixelmatch accurately compares screenshots by analyzing color differences, easily distinguishing between intentional brand color updates and broken page layouts. Tests are validated across multiple device viewports before any change is committed.

TESTING FIXES AUTOMATICALLY BEFORE COMMIT: When the auto-healer writes a new selector, it tests the fix inside a headless browser across mobile, tablet, and desktop screens. Only when the test passes on all screens does it commit the fix, saving developers from boring maintenance chores while providing clear visual diff snapshots on every pull request.

SAVING SPRINT TIME FOR PRODUCTIVE WORK: Every hour a developer spends updating brittle test selectors is an hour lost from building new product features. By healing 89.4% of broken locators automatically and providing clear before-and-after visual diffs, MACT removes testing frustration and lets software teams focus on what matters most.

7. DUAL-MODEL SEMANTIC ROUTER: SMART AI DISPATCHING

Sending simple tasks to a fast lightweight AI model and complex tasks to a deep reasoning model cuts AI costs by 68% and speeds up response times.




Figure 5. Dual-Model Semantic Router Architecture directing simple checks to the Fast Path and hard tasks to the Deep Reasoning Path. [4]

HOW THE AI ROUTER DISPATCHES TASKS

Our dispatch topology routes verification tasks with high efficiency. Incoming verification tasks are checked by the Semantic Router. Simple tasks (like formatting checks, quick syntax linting, and basic unit tests) are routed to the Fast Path (GPT-5.6 Sol). Complex tasks (like security data flow analysis, multi-service debugging, and UI self-healing) are sent to the Deep Reasoning Path (Claude-Opus-5). In-memory cached symbol tables speed up prompts, and a response gate ensures answers follow clean schemas.

SPEED AND COST EFFICIENCY

Sending every test task to an expensive frontier model is unnecessary and wasteful. Simple code checks do not need deep multi-step reasoning. By sending 78% of routine checks to the fast model and saving the deep reasoning model for hard problems, MACT lowers verification costs to just \$1.54 per pull request while keeping routing latency under 300 milliseconds.

BALANCING SPEED, COST, AND ANALYTICAL ACCURACY

The semantic router calculates task difficulty based on how many files and functions were changed. Low-complexity tasks finish in milliseconds on the fast path, keeping costly reasoning budgets available for deep architectural checks.

Pre-cached code symbols in local server memory cut prompt sizes by 74%, allowing teams to run parallel test sweeps across dozens of pull requests without hitting AI provider rate limits.

SMART BUDGETING & PREDICTABLE COSTS: Both models output structured JSON validated against strict rules before results are posted to the shared blackboard. If a model returns invalid text, the system asks it again with zero randomness, ensuring tests always produce reliable, repeatable results at a fraction of the cost.

DYNAMIC ROUTING BASED ON REAL CODE CHANGES: If a pull request only touches documentation or simple formatting, the semantic router handles the entire verification on the fast path in less than 2 seconds. When complex business logic or multi-service APIs change, the router automatically pulls in the deep reasoning model. This adaptive intelligence ensures optimal speed for small commits and maximum depth for complex changes.

8. PULL REQUEST TO VERDICT: FAST EXECUTION SEQUENCE

A clear step-by-step timeline shows how code changes move from a Git pull request to a signed approval verdict in under 15 seconds.




Figure 6. Pull Request to Release Verdict Sequence Flow showing fast, parallel communication across all 6 system components. [5]

STEP-BY-STEP WORKFLOW TIMELINE

Our execution sequence maps the complete journey of a code verification run across 6 key components: Developer/PR, Master Orchestrator, TupleGrid Blackboard, Cognitive Agents Grid, MCP Tool Boundaries, and Consensus Reporter. (1) A pull request triggers a webhook. (2) The orchestrator publishes task tickets to the blackboard. (3) Specialized agents wake up simultaneously. (4-5) Agents run tools inside safe MCP containers and collect results. (6) Findings are posted back to the board. (7-8) The arbiter reviews all rules and publishes the final release decision.

SUB-15s SPEED

Multi-agent verification achieves enterprise speed without cutting corners. By replacing sequential linear pipeline steps with non-blocking blackboard leasing, MACT completes comprehensive verification sweeps in under 15 seconds. Independent agent execution prevents single-tool stalls from blocking unrelated domains, maintaining predictable developer feedback velocity.

PRE-WARMED WORKERS & NON-BLOCKING TASK LEASES

Measuring the full sequence shows how fast each step runs: Reading changes takes 0.4s, publishing task plans takes 0.8s, running tools across all four pillars in parallel takes 11.8s, gathering results takes 0.6s, and making the final release decision takes 0.9s—totaling just 14.5 seconds.

Because tools run in pre-warmed container pools with ready runtimes (Node.js, Go, Python), startup delays are completely eliminated. Tasks expire automatically if an agent fails to respond in time, allowing the watchdog to reassign tasks without hanging the pipeline.

HOW NON-BLOCKING STEPS DROP WAIT TIMES: Old-fashioned CI pipelines run tests in a straight line: unit tests first, then integration tests, then security scans, then UI tests. MACT puts all tasks on the shared TupleGrid board at once. All agents grab their work simultaneously and run tools in parallel, finishing everything in just 14.5 seconds.

HIGH-VELOCITY FEEDBACK KEEPS DEVELOPERS IN FOCUS: Waiting 45 minutes for a build breaks engineering momentum. Getting test results in under 15 seconds means developers can fix any small typo or broken test immediately while the code is still fresh in their mind, radically increasing sprint delivery velocity.

9. TUPLEGRID IN-MEMORY BLACKBOARD COORDINATION

A fast in-memory blackboard lets agents coordinate without direct messaging, avoiding network bottlenecks and preventing communication deadlocks.

9.1 Why Direct Agent-to-Agent Messaging Fails

When many AI agents try to message each other directly, the system quickly slows down. With 19 domain agents operating across all verification modules, direct messages multiply rapidly, causing network traffic jams, high serialization overhead, and deadlocks. If one agent pauses or crashes while waiting for a response from another, the entire testing pipeline grinds to a halt.

MACT solves this by using **TupleGrid**: an in-memory shared blackboard. Agents never send messages directly to each other. Instead, they share work by reading and posting clear data records to a shared memory board, completely avoiding communication gridlock.

9.2 The Four Core Coordination Records

TupleGrid organizes work using four simple, strongly-typed memory records:

#	Record Name	Key Information Included	How It Is Used in the Workflow
1	DIFF_DAG_PLAN	pr_id, commit_sha, task_list, timeout_ms	Published by Orchestrator; defines tasks and dependencies. Tasks expire automatically if an agent gets stuck.
2	AGENT_STATUS	agent_id, pillar_id, current_state, heartbeat	Tracks agent state (IDLE / WORKING / FINISHED). The watchdog monitors heartbeats to reassign timed-out tasks.
3	DOMAIN_FINDING	pillar_id, module_id, test_result, audit_hash	Permanent test findings posted by agents. Cryptographic hashes guarantee that results cannot be tampered with.
4	MCP_TOOL_RESULT	tool_id, exit_code, output_text, duration_ms	Standard execution result returned from isolated tool containers, providing concrete diagnostic proof.

Table 3. TupleGrid In-Memory Coordination Records & Rules [4]

FAST LOCK-FREE COORDINATION & HIGH-THROUGHPUT RELIABILITY

To eliminate waiting, TupleGrid uses atomic memory operations. When an agent wants to claim a task, it updates the record instantly in memory. If two agents try to grab the same task at the exact same moment, the collision is resolved in under 20 microseconds without slowing down other workers.

Agents subscribe to the specific test findings they care about. As soon as a matching record is posted to the board, subscribed agents wake up instantly with zero delay.

WHY A SHARED BLACKBOARD STOPS COMMUNICATION DEADLOCKS: Because all temporary coordination happens in fast memory without waiting on disk writes, TupleGrid processes over 18,000 operations per second with response times under 2.2ms, keeping multi-agent teamwork smooth and responsive under heavy workloads.

INSTANT RECOVERY IF AN AGENT CRASHES: Every task posted to the board has an automatic heartbeat. If an agent crashes or loses connection while running a tool, the watchdog notices the expired timer and hands the task to a healthy agent within 5 seconds, ensuring builds never stall.

RELIABLE AUDIT TRAIL WITHOUT DISK OVERHEAD: All verification findings, execution durations, and tool outputs stay in memory throughout the pull request lifecycle. When the release decision is made, the full record is signed and saved, giving teams a complete audit history of every test without slowing down active testing runs.

DECOUPLING WORKERS FROM TIMING DEPENDENCIES: Unlike message queues where producers and consumers must be closely coordinated, TupleGrid decouples agents in both time and space. A fast agent can post its results and immediately claim the next task without waiting for downstream reviewers to come online, maximizing server utilization.

10. DYNAMIC CONSENSUS DECISION ENGINE

A smart decision tree combines findings from all four pillars into a single release verdict, making sure critical security issues always stop a release.




Figure 7. Dynamic Invariant Consensus Decision Tree showing how findings from all 19 modules are evaluated to approve or block a release. [6]

HOW THE CONSENSUS ENGINE DECIDES

Our decision topology evaluates test results from all four pillars through a centralized consensus engine: Functional (Pillar 1), Security (Pillar 3), Performance (Pillar 2), and Quality (Pillar 4). The Consensus Arbitrator checks these findings against strict priority rules. If all required checks pass, the release is marked **VERIFIED STABLE** and allowed to merge. If any critical rule is broken (like a security vulnerability or broken API), release is **BLOCKED** immediately.

ZERO UNCHECKED DEFECTS

Software release decisions must follow strict priority rules rather than simple majority voting. In majority voting, if five tests pass but one security scan finds a serious SQL injection bug, a 5-to-1 vote would dangerously approve the release! MACT's rule hierarchy ensures that security always has absolute veto power: a single security flaw stops the release immediately.

HOW CONFLICTING TEST RESULTS ARE RESOLVED VIA COGNITIVE ARBITRATION

When Pillar 1 (Functional) passes but Pillar 2 (Performance) detects that the application has become clearly slower, the arbitrator resolves the conflict using strict priority principles: functional correctness cannot excuse performance degradation. Instead of relying on a rigid, hardcoded threshold (such as an arbitrary 15% rule), the performance agent queries historical benchmark baselines and latency distributions persisted in the database (`mact_analytics_db`). The LLM agent evaluates current telemetry against this historical database context, using cognitive reasoning to determine whether a measured slowdown (whether 12%, 15%, 20%, or more) constitutes an unacceptable regression for that workload. If confirmed, the release is paused, historical baseline comparison logs and diagnostic traces are attached, and engineers are notified.

The priority hierarchy is clear and unchanging: Security & Compliance (Priority 1) > Functional Contracts & UI (Priority 2) > Performance (Priority 3) > Code Quality Policies (Priority 4). This guarantees that vulnerable, functionally broken, or severely degraded code cannot slip into production.

WHY SECURITY ALWAYS OVERRULES REGULAR TESTS: In software development, teams are sometimes tempted to ignore security warnings if feature tests pass. MACT eliminates human bias by enforcing an absolute rule: Security is Priority 1. Every approved release includes a digitally signed summary of all test results, providing verified confidence in every deployed build.

COGNITIVE LLM REASONING REMOVES SUBJECTIVITY: Traditional CI systems rely on naive static thresholds that cause false alarms or miss systemic regressions. MACT's cognitive agents evaluate code diffs against historical database memory, ensuring objective, dependable release standards without manual arguments.

11. KERNEL OBSERVABILITY & eBPF DIAGNOSTIC TRIAGE

Using lightweight eBPF kernel probes uncovers the true cause of test timeouts, eliminating 98.4% of flaky test false alarms and synthesizing automatic code fixes.




Figure 8. eBPF Observability & Automated Healing Loop diagnosing root causes and creating verified pull request fixes.
[7]

THE CLOSED-LOOP DIAGNOSTIC TRIAGE

Our kernel observability topology operates across 6 distinct phases: (1) A test fails unexpectedly. (2) The diagnostic engine checks eBPF kernel network socket events and CPU timings. (3) The classifier determines whether the failure was caused by real code errors or just temporary server lag. (4) If it is a real code bug, the synthesizer drafts a targeted code fix. (5) The fix is verified in a clean temporary container. (6) Once proven green, the fix is committed to the pull request.

STOPPING FLAKY TEST FALSE ALARMS

Looking directly into system-level kernel events solves the flaky test problem once and for all. By checking whether server CPU was overloaded or network buffers were full, MACT objectively proves whether a timeout was caused by slow infrastructure or a genuine code bug. This eliminates 98.4% of false alarms, preventing teams from chasing phantom bugs.

ACCURATE ROOT-CAUSE ISOLATION & TESTED PATCHES

If a generated code patch fails its verification test, the feedback loop immediately sends the failure log back to the synthesizer with clear instructions on what did not work. This feedback ensures that only fully working, verified fixes are ever committed to your repository.

eBPF probes run with virtually zero performance impact (under 0.5% CPU overhead). If a test times out while the server CPU was heavily congested by unrelated jobs, the failure is flagged as server congestion and scheduled for a clean re-run instead of failing the developer's build.

CLOSED-LOOP LEARNING TO FIX CODE SAFELY: When a real bug is found and a patch is generated, MACT tests the patch in a clean container. If the patch fails, the system learns what went wrong and tries a better fix, ensuring only working code is ever committed.

WHY eBPF PROBES ARE ESSENTIAL FOR CI CLUSTERS: In busy Kubernetes clusters, multiple test suites share the same underlying servers. When several builds start at once, CPU and network resources spike temporarily. Traditional test runners see this as a test failure. eBPF probes inspect kernel socket queues directly, proving that the code was fine and the server was just busy. This simple insight stops teams from wasting time on false alarms.

12. THE COMPLETE UNIFIED SYSTEM TOPOLOGY CANVAS

The complete 5-layer system overview shows how triggers, AI orchestration, shared memory, agent teams, and secure tools operate as one unified platform.




Figure 9. Multi-Agent Orchestration Canvas showing the 5 connected layers from pull request trigger down to isolated tool execution. [8]

THE 5-LAYER SYSTEM TOPOLOGY

Our unified system topology provides the complete architectural picture of the MACT platform across 5 connected layers: Layer 1 (Ingestion & Trigger: reads Git webhooks and analyzes changed code files); Layer 2 (Cognitive Orchestration: manages task planning and picks the best AI model); Layer 3 (TupleGrid Blackboard: manages shared in-memory task coordinates); Layer 4 (Specialized Agents: 19 focused domain agents across the 4 pillars, 21 platform agents total); and Layer 5 (MCP Tool Suites: all 19 sidebar module tools running safely inside isolated containers).

COMPLETE SYSTEM SYNERGY

This unified topology demonstrates that autonomous testing requires thoughtful structure across all layers. Smart agents without safe containers would create security risks; isolated tools without an in-memory board would cause messaging traffic jams; and testing without smart routing would cost too much. Connecting all five layers into one smooth system gives MACT its speed, safety, and intelligence.

HOW WORK FLOWS SEAMLESSLY FROM TOP TO BOTTOM

Instructions flow cleanly from top to bottom, while test results and diagnostic traces flow back up through verified memory records. This clear discipline prevents unauthorized actions while ensuring that release decisions are grounded in real, verifiable test evidence.

The shared memory board at Layer 3 completely decouples planning from execution: orchestrators can plan tasks without needing to know tool syntax, while tool runners execute jobs without needing to know prompt details. This clean separation makes MACT easy to maintain and expand.

WHY A 5-LAYER DESIGN MAKES THE SYSTEM DURABLE: In poorly designed systems, tool scripts and AI prompts are tangled together. In MACT, you can upgrade any tool in Layer 5 or swap any AI model in Layer 2 without touching the rest of the platform, making MACT built for long-term maintainability.

CLEAN STRATIFIED TOPOLOGY FOR MODERN ENTERPRISES: The 5-layer design guarantees that every layer has a clear, isolated responsibility. From the moment code is pushed to Git down to the sandboxed execution of testing tools, every step is predictable, auditable, and resilient to failures, creating a modern standard for autonomous software engineering.

13. HYBRID MULTI-AGENT PATTERNS & COORDINATION TOPOLOGY

MACT integrates four foundational distributed systems patterns into a unified hybrid multi-agent topology: Hierarchical Orchestration, Space-Based Blackboard Coordination, DAG Pipeline Staging, and Strict Consensus Arbitration.




Figure 10. Hybrid Multi-Agent Architectural Topology Canvas (Archify Showcase Architecture): Stratified coordination from Git ingestion down through TupleGrid in-memory space, 4-pillar agent teams, and consensus release gating. [10]

THE FOUR FOUNDATIONAL HYBRID PATTERNS

MACT combines four distinct multi-agent patterns into a cohesive engine: 1) Hierarchical Orchestrator-Worker: Master agent `agent_coordinator` plans global DAG execution without prompt drift while 19 domain agents execute specialized verifications; 2) Space-Based Blackboard (TupleGrid): In-memory tuples decouple producers from consumers, eliminating severe point-to-point network bottlenecks; 3) DAG Pipeline: Slices AST diffs to execute fast static checks upstream before dynamic tests; and 4) Consensus Arbiter: Synthesizes cross-pillar invariants with strict safety-first priority (Security > UX > Perf > IaC).

WHY MACT REJECTED PEER-TO-PEER MESH & SWARMS

Traditional multi-agent frameworks often adopt peer-to-peer mesh or autonomous swarm patterns. In an enterprise verification platform with 19 domain agents, point-to-point mesh communication creates dozens of chaotic interconnected channels, causing message saturation, high token costs, and circular wait deadlocks. Pure swarms lack predictable auditability required for compliance. MACT's hybrid blackboard model guarantees sub-2.2ms coordination latency with total reliability.

PRODUCTION PATTERN TAXONOMY & ENTERPRISE TRADE-OFFS

In distributed multi-agent systems, pattern selection dictates resilience and failure modes. Pure swarms and mesh networks rely on gossip protocols and direct agent-to-agent negotiations. When an agent experiences transient network latency or LLM rate limits, point-to-point topologies suffer cascading timeouts. MACT eliminates this failure mode entirely through spatial memory decoupling: agents only interact via TupleGrid leases, ensuring no agent ever blocks waiting on another.

PREDICTABLE AND BOUNDED COORDINATION: By enforcing a clean hierarchical DAG structure, MACT focuses strictly on modified code branches and directly affected test suites. This eliminates runaway coordination loops and guarantees sub-1.2s dispatch times regardless of repository scale, providing predictable, auditable continuous verification for enterprise environments.

STRICT INVARIANT PRIORITY IN PRODUCTION RELEASES: Consensus is not reached by majority vote or unweighted averaging. MACT enforces an immutable safety precedence: *Security Invariants* > *UX Functional Correctness* > *Latency/SRE Budgets* > *Cloud IaC Contracts*. If Pillar 3 discovers a critical taint injection or unpatched CVE, the release is blocked unconditionally, regardless of flawless 100% test passes across the remaining 14 agents. This deterministic gate guarantees zero-defect releases.

14. THE AUTONOMOUS AGENT VERIFICATION HARNESS (<https://harness.mact.pro>)

Deterministic Micro-Sandboxes, Continuous Verification (CV) Gates, and Zero-Trust Tool Isolation for Multi-Agent Workflows.




Figure 11. Autonomous Agent Verification Harness (<https://harness.mact.pro>) Architecture & Sandbox Lifecycle: Deterministic micro-sandboxes, sub-40ms warm pool, boundary invariant fuzzing, and analytics-driven continuous verification gates. [11]

THE 4-TIER HARNESS EXECUTION & CV LIFECYCLE

Our verification harness (<https://harness.mact.pro>) separates high-level cognitive reasoning from low-level execution across 4 distinct stages: (1) Cognitive agents dispatch signed JSON-RPC execution tickets. (2) Container A (Security Gateway) checks parameters against strict JSON schemas and leases scoped capability tokens. (3) Container B (Harness Sandbox Grid) boots an ephemeral micro-sandbox in <38ms with read-only root filesystems and isolated memory. (4) Continuous Verification (CV) asserts operational metrics & performance compliance (Prometheus, Jaeger) before the Release Arbiter issues a signed verdict.

ZERO-DRIFT ISOLATION & HOST PROTECTION

Direct shell execution by autonomous LLM agents presents catastrophic security and reliability risks. An unchecked prompt can execute destructive commands, pollute shared test environments, or leak cloud credentials. Harness eliminates this vulnerability through zero-drift, disposable microVM sandboxes: every test suite runs in an isolated RAM disk that is completely destroyed upon test completion, guaranteeing 100% reproducible results and zero environment leakage.

DETERMINISTIC INVARIANTS & CANARY GATES

Cognitive agents such as agent_fuzz and agent_mutation dynamically inject boundary invariants, synthesized property tests, and semantic mutants into the harness sandbox without exposing the host system. By evaluating function signatures against thousands of edge-case variations in parallel, Harness proves code robustness before merge.

Beyond binary pass/fail assertions, the Continuous Verification (CV) engine continuously monitors live performance analytics. If a pull request introduces an abnormal memory consumption trend (>0.00% leak drift) or breaches the P99 latency threshold (<=156ms), Harness halts deployment, triggers automated canary rollback, and attaches diagnostic eBPF traces for developer remediation.

WHY EPHEMERAL MICRO-SANDBOXES DEFEAT CI BOTTLENECKS: Traditional CI runners require 30 to 60 seconds of cold setup time to pull containers and initialize dependencies. Harness maintains a pre-warmed pool of micro-sandboxes with ready runtimes (Node.js, Go, Python), dropping initialization time to under 38 milliseconds and accelerating verification sweeps to under 15 seconds.

ENTERPRISE CONTINUOUS VERIFICATION AT SCALE: In production environments, passing unit tests does not guarantee operational safety. By combining MCP capability leases, isolated boundary fuzzing, and analytics-correlated release gates, Harness establishes an autonomous, zero-trust verification foundation for enterprise multi-agent engineering.